

CS 1550 Project 3 Write Up

Dr. Mosse - Spring 2019
Mathew Varughese
mav120@pitt.edu

Aging Algorithm

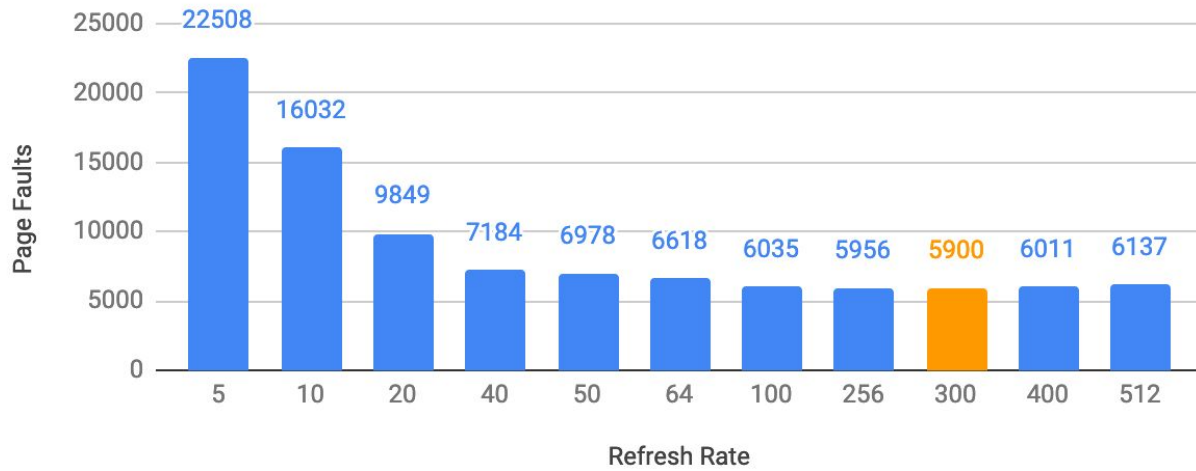
The aging algorithm takes in a “refresh” parameter. This refresh parameter is used to determine how many clock cycles should happen between each refresh. When a refresh occurs, all of the pages in the frame table have their aging counter shifted right by one. Effectively, this refresh ages the pages that are not often used. However, determining an appropriate refresh parameter is tricky. It turns out to be quite dependent on the specific file and the number of frames used. To determine a parameter, I first used 16 frames and tested the algorithm on different refresh rates.

This data is shown in this table below:

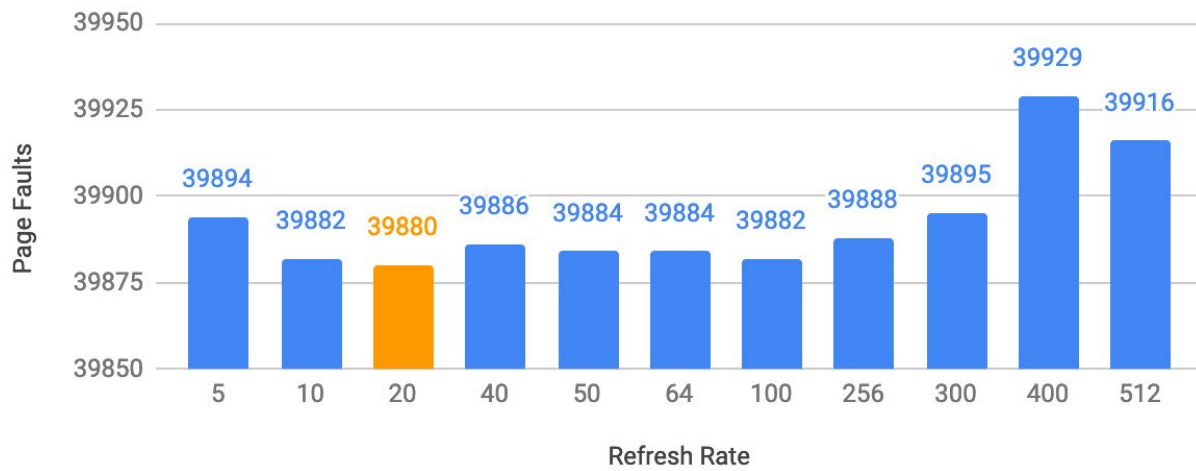
gcc.trace		gzip.trace		swim.trace	
Refresh Rate	Page Faults	Refresh Rate	Page Faults	Refresh Rate	Page Faults
5	22508	5	39894	5	21001
10	16032	10	39882	10	5328
20	9849	20	39880	20	997
40	7184	40	39886	40	752
50	6978	50	39884	50	689
64	6618	64	39884	64	601
100	6035	100	39882	100	585
256	5956	256	39888	256	560
300	5900	300	39895	300	559
400	6011	400	39929	400	569
512	6137	512	39916	512	579

Graphing this data helped give a little insight into which parameter works the best. These can be seen on the next page. However, it does not give a clear picture which refresh parameter to choose. For `gcc` and `swim`, a parameter of 300 seemed to work. For `gzip`, the refresh rate parameter is 20. This is because it seems dependent on the number of frames used. On the page after this there is a chart that compares refresh rate and different frames.

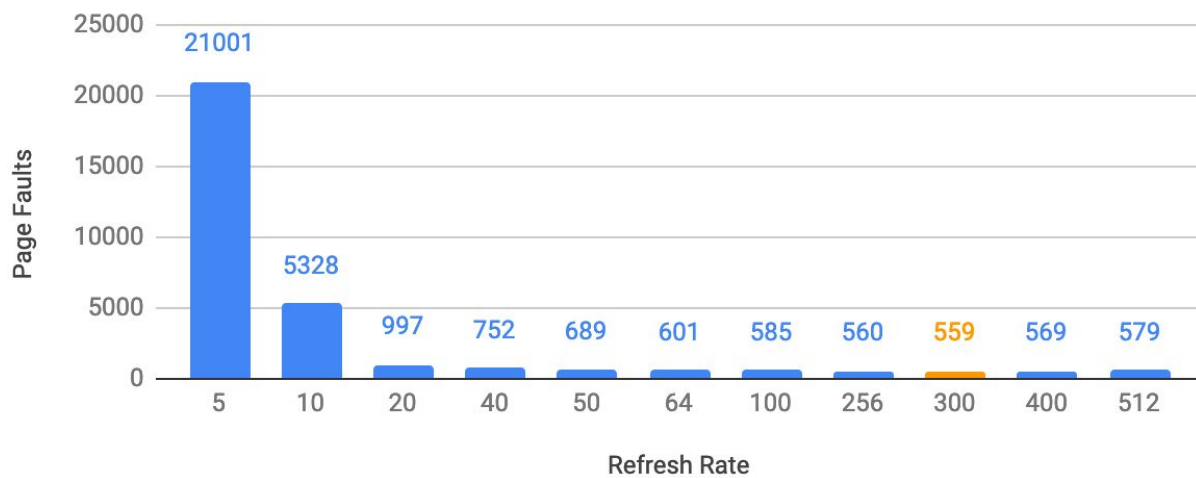
AGING - gcc.trace - (16 Frames)



AGING - gzip.trace - (16 Frames)



AGING - gcc.trace - (16 Frames)



Refresh Rate vs Page Faults per Frame				
gcc.trace	Page Faults			
Refresh Rate	8 Frames	16 Frames	32 Frames	64 Frames
16	22400	13471	12510	6474
32	22413	7454	5588	3266
64	22771	6618	3194	1910
128	23436	5976	1732	1213
256	24499	5956	1048	795
400	25268	6011	837	607
512	25801	6137	783	566
650	27101	6173	761	553
gzip.trace	Page Faults			
Refresh Rate	8 Frames	16 Frames	32 Frames	64 Frames
16	39903	39879	39875	39875
32	39902	39882	39874	39874
64	39916	39884	39874	39874
128	39998	39883	39874	39872
256	40089	39888	39874	39874
400	40049	39929	39875	39874
512	40203	39916	39874	39874
650	40462	40141	39890	39874
swim.trace	Page Faults			
Refresh Rate	8 Frames	16 Frames	32 Frames	64 Frames
16	9297	2115	1286	361
32	8127	825	618	319
64	6595	601	444	301
128	6377	575	376	290
256	5885	560	277	232
400	6468	569	231	198
512	6684	579	213	185
650	6861	585	208	174

The graphs for this data in shown on the next page. From here, it seems that a refresh rate of about 256 seems to be sufficient for the purposes of this project. Although intuitively it seems like a small refresh rate would be better, more time allows more pages that will be accessed to turn their reference bit on. It would be better to have more traces to test to be more certain of an answer.


Comparing Page Replacement Algorithms

The table below shows the page faults and disk writes of performing the simulation on different frames and using different algorithms.


gcc.trace	Page Faults			Disk Writes		
Frames	OPT	FIFO	AGING	OPT	FIFO	AGING
8	4380	11519	24499	13328	29011	24499
16	1019	3542	5956	3020	8568	5956
32	215	660	1048	491	1375	1048
64	105	258	795	318	551	795
swim.trace	Page Faults			Disk Writes		
Frames	OPT	FIFO	AGING	OPT	FIFO	AGING
8	2182	8499	5885	4417	13893	5885
16	149	470	560	358	844	560
32	60	158	277	144	326	277
64	26	69	232	135	177	232
gzip.trace	Page Faults			Disk Writes		
Frames	OPT	FIFO	AGING	OPT	FIFO	AGING
8	39844	39894	40089	39874	44918	40089
16	39825	39856	39888	39856	42384	39888
32	39809	39825	39874	39856	41120	39874
64	39777	39793	39874	39856	40496	39874

Charts on the next pages display a better picture of this data.


Page Faults - gcc.trace




Page Faults - swim.trace



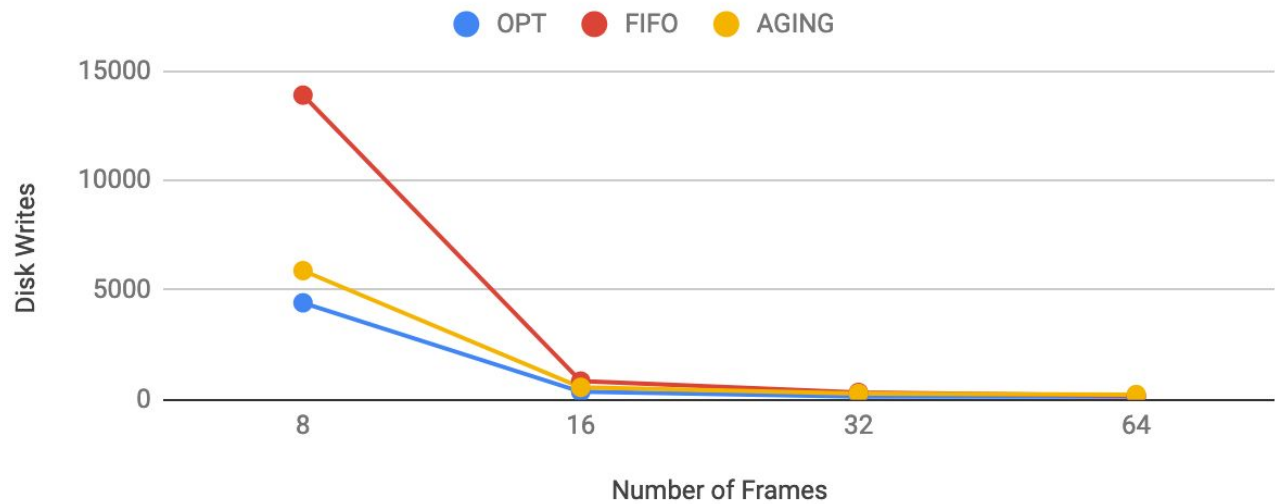
Page Faults - gzip.trace




Disk Writes - gcc.trace



Disk Writes - swim.trace



Disk Writes - gzip.trace



The data shows that increasing the number of page frames decreases the number of page faults and disk writes. This is as expected. Other than cases of Belady's Anomaly, which will be explained in the next section, increasing the number of frames will increase the number of pages stored. This means page faults occur less because there is a much higher probability a page that needs to be found can be in the page table.

These charts also depict the idea of the law of diminishing returns. The law of diminishing returns is an economic concept that shows that not every increase in input will give a linear increase in output. At a point, after adding more input, there is decreasing benefit to the rate of return. After analyzing the number of frames, it is evident that increasing them from 8 to 16 gives a substantially lower number of page faults. The difference between 32 and 64 is not so large however. This is because as the number of frames grows to infinity, the number of page faults decreases. However, this might be unnecessary. Not all programs need to use an infinite number of frames.

This is where the importance of picking an appropriate page replacement algorithm comes into play. A smart page replacement algorithm is able to decide which frames to evict with higher confidence. This in turn lowers the number of page faults. The optimal page replacement algorithm can only occur in a perfect world where the Operating System somehow knows exactly which memory calls will be made when. It chooses which page to evict based on which one will be used the farthest into the future. In the graphs, it is evident the optimal page replacement algorithm performs the best. It can be mathematically shown that it is better than the other two.

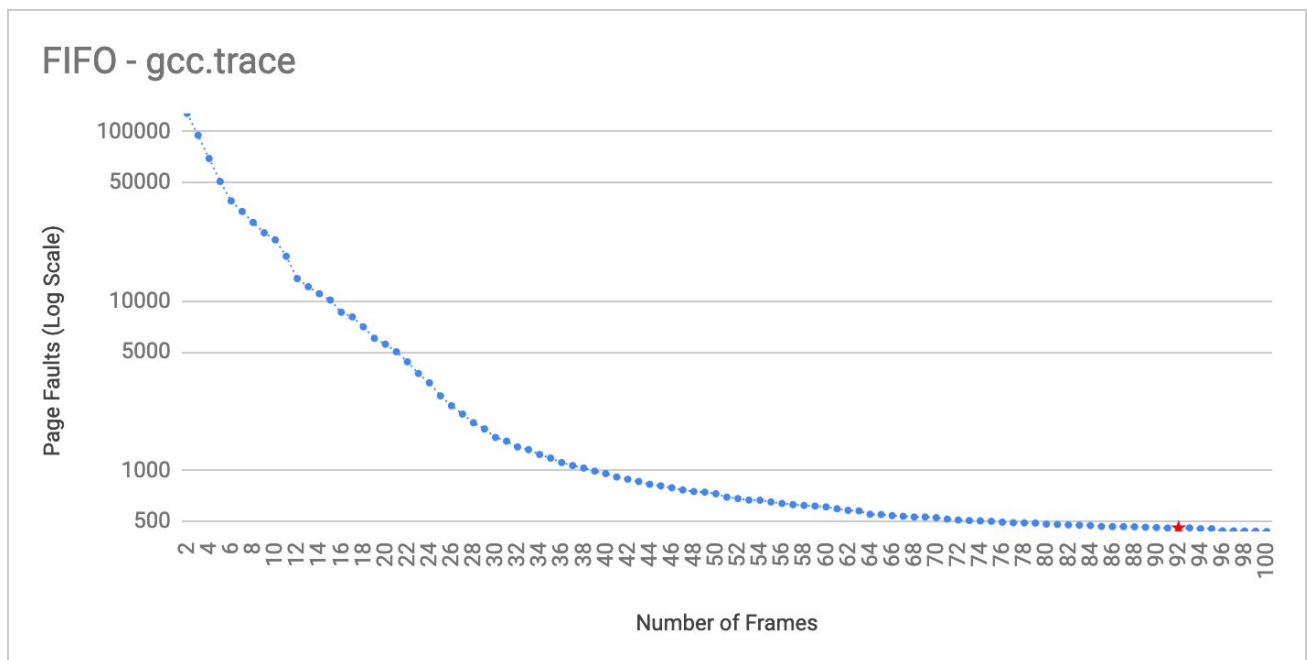


First-In-First-Out is a simple algorithm that does not perform that great. It decides to evict the first page that was inserted into the page table. In the chart, it surprisingly performed better than the aging algorithm in some cases. For example, in `gzip`, it was able to do significantly less page faults than the aging algorithm. Again, this is program dependent. Since `gzip` is a compression program, it likely does not use the same memory locations often. However, this is speculation. Based on the other programs, FIFO performs more disk operations and usually more page faults. Writing pages back to disk occur when a page is dirty. This is extremely slow and should be avoided at all costs.

For this reason, the *aging algorithm* is likely the best page replacement algorithm to use in an actual Operating System. Of course, this is dependent based on the data and programs given. The aging algorithm could adjust the refresh rate based on machine learning or other algorithms to determine the best refresh rate for each process. 256 was used as a rough estimate as optimal here, but it is hard to tell for sure. Machine learning would be extremely useful here in using past data to predict future behavior in order to pick the most efficient values.

Therefore, I believe the aging algorithm is most appropriate to use in an actual operating system. It performs disk writes at a rate much closer to the optimal algorithm than FIFO which make it a strong candidate. Even though it causes more page faults than the optimal algorithm, this could be remedied through the use of machine learning algorithms to adjust the refresh parameter.

Belady's Anomaly

Belady's Anomaly is a phenomenon where increasing the number of page frames results in an increase in the number of page faults for a given memory access pattern. This is very counterintuitive. However, it occurs in the `gcc` program. To find this, here are three graphs that show the number of page faults that occur as the number of frames grows from 2 to 10. Belady's Anomaly occurs when the number of frames is equal to 92 on `gcc.trace`. The table data is shown on the next page.



Frames (gcc.trace)	Page Faults	Difference		51	695	33
2	127688	0		52	682	13
3	94726	32962		53	668	14
4	69158	25568		54	667	1
5	50623	18535		55	651	16
6	38848	11775		56	638	13
7	33664	5184		57	628	10
8	29011	4653		58	621	7
9	25202	3809		59	616	5
10	22916	2286		60	609	7
11	18325	4591		61	594	15
12	13529	4796		62	581	13
13	12118	1411		63	577	4
14	11014	1104		64	551	26
15	10111	903		65	549	2
16	8568	1543		66	541	8
17	8045	523		67	536	5
18	7026	1019		68	531	5
19	6018	1008		69	531	0
20	5550	468		70	526	5
21	5011	539		71	516	10
22	4370	641		72	510	6
23	3728	642		73	506	4
24	3283	445		74	504	2
25	2752	531		75	501	3
26	2407	345		76	495	6
27	2148	259		77	491	4
28	1910	238		78	490	1
29	1755	155		79	489	1
30	1565	190		80	482	7
31	1488	77		81	480	2
32	1375	113		82	477	3
33	1325	50		83	475	2
34	1241	84		84	473	2
35	1182	59		85	467	6
36	1112	70		86	466	1
37	1068	44		87	466	0
38	1032	36		88	465	1
39	990	42		89	462	3
40	958	32		90	460	2
41	914	44		91	457	3
42	887	27		92	458	-1
43	860	27		93	458	0
44	828	32		94	453	5
45	811	17		95	453	0
46	790	21		96	440	13
47	766	24		97	440	0
48	751	15		98	439	1
49	744	7		99	438	1
50	728	16		100	436	2